




# Intro to ML

~~Live 08~~ - Data Science Bootcamp  
Live 10

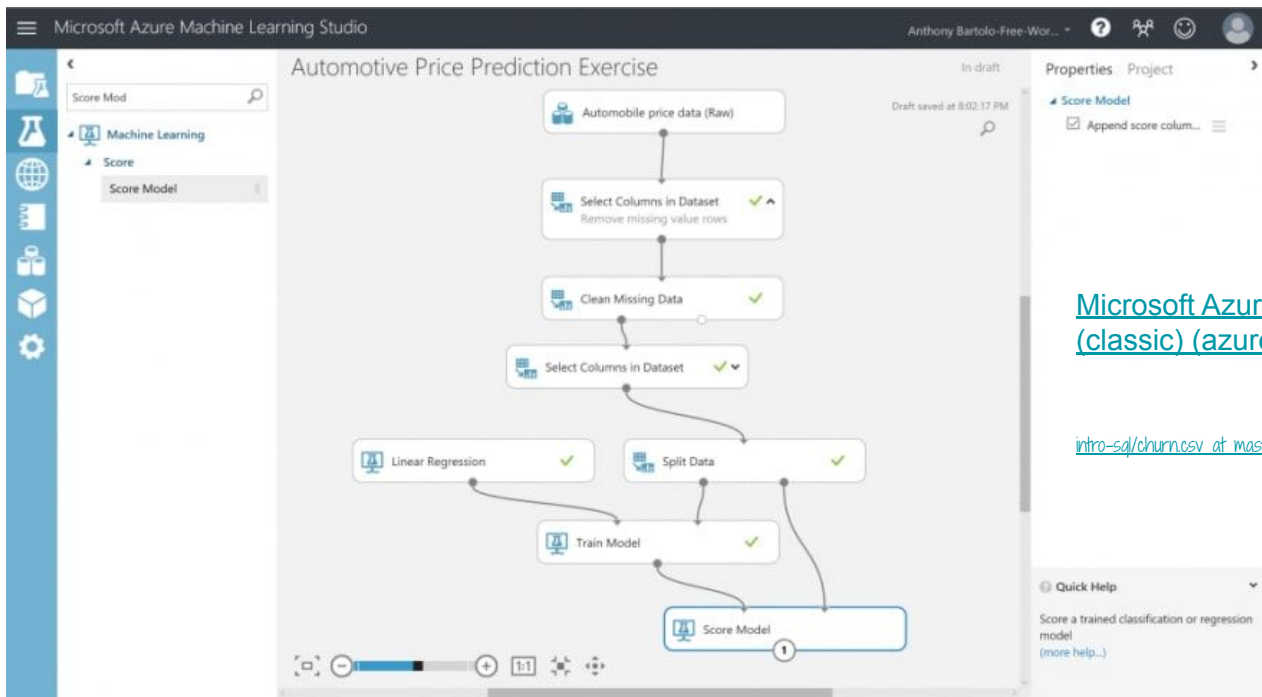


## Simple pipeline to build ML models





# Build your first model with Azure ML Studio



[Microsoft Azure Machine Learning Studio \(classic\) \(azureml.net\)](https://azureml.net)

[intro-sql/churn.csv at master · toyeyei/intro-sql \(github.com\)](https://github.com/toyeyei/intro-sql)

Example dataset

# Machine Learning

When a computer can learn to recognize pattern



## **Essential ML**

- what exactly is machine learning
- supervised vs. unsupervised
- regression vs. classification
- train test split vs. cross validation
- model selection + hyperparameter
- model evaluation



R

# What is Machine Learning



Arthur Samuel (1959)


Field of study that gives computers **the ability to learn without being explicitly programmed.**

## R How do we learn?



Human learn from **experience**.  
Computer learn from **data**.

**R Simple Idea**







## ML Glossary #1

- dataset
- data points
- features
- label or target

Data Point

	crim	zn	indus	chas	nox	rm	age	dis	rad	tax	ptratio	b	lstat	medv
1	0.00632	18.0	2.31	0	0.5380	6.575	65.2	4.0900	1	296	15.3	396.90	4.98	24.0
2	0.02731	0.0	7.07	0	0.4690	6.421	78.9	4.9671	2	242	17.8	396.90	9.14	21.6
3	0.02729	0.0	7.07	0	0.4690	7.185	61.1	4.9671	2	242	17.8	392.83	4.03	34.7
4	0.03237	0.0	2.18	0	0.4580	6.998	45.8	6.0622	3	222	18.7	394.63	2.94	33.4
5	0.06905	0.0	2.18	0	0.4580	7.147	54.2	6.0622	3	222	18.7	396.90	5.33	36.2
6	0.02985	0.0	2.18	0	0.4580	6.430	58.7	6.0622	3	222	18.7	394.12	5.21	28.7
7	0.08829	12.5	7.87	0	0.5240	6.012	66.6	5.5605	5	311	15.2	395.60	12.43	22.9
8	0.14455	12.5	7.87	0	0.5240	6.172	96.1	5.9505	5	311	15.2	396.90	19.15	27.1
9	0.21124	12.5	7.87	0	0.5240	5.631	100.0	6.0821	5	311	15.2	386.63	29.93	16.5
10	0.17004	12.5	7.87	0	0.5240	6.004	85.9	6.5921	5	311	15.2	386.71	17.10	18.9
11	0.22489	12.5	7.87	0	0.5240	6.377	94.3	6.3467	5	311	15.2	392.52	20.45	15.0
12	0.11747	12.5	7.87	0	0.5240	6.009	82.9	6.2267	5	311	15.2	396.90	13.27	18.9
13	0.09378	12.5	7.87	0	0.5240	5.889	39.0	5.4509	5	311	15.2	390.50	15.71	21.7
14	0.62976	0.0	8.14	0	0.5380	5.949	61.8	4.7075	4	307	21.0	396.90	8.26	20.4
15	0.63796	0.0	8.14	0	0.5380	6.096	84.5	4.4619	4	307	21.0	380.02	10.26	18.2

Dataset: BostonHousing

## Features (X)

	crim	zn	indus	chas	nox	rm	age	dis	rad	tax	ptratio	b	lstat	medv
1	0.00632	18.0	2.31	0	0.5380	6.575	65.2	4.0900	1	296	15.3	396.90	4.98	24.0
2	0.02731	0.0	7.07	0	0.4690	6.421	78.9	4.9671	2	242	17.8	396.90	9.14	21.6
3	0.02729	0.0	7.07	0	0.4690	7.185	61.1	4.9671	2	242	17.8	392.83	4.03	34.7
4	0.03237	0.0	2.18	0	0.4580	6.998	45.8	6.0622	3	222	18.7	394.63	2.94	33.4
5	0.06905	0.0	2.18	0	0.4580	7.147	54.2	6.0622	3	222	18.7	396.90	5.33	36.2
6	0.02985	0.0	2.18	0	0.4580	6.430	58.7	6.0622	3	222	18.7	394.12	5.21	28.7
7	0.08829	12.5	7.87	0	0.5240	6.012	66.6	5.5605	5	311	15.2	395.60	12.43	22.9
8	0.14455	12.5	7.87	0	0.5240	6.172	96.1	5.9505	5	311	15.2	396.90	19.15	27.1
9	0.21124	12.5	7.87	0	0.5240	5.631	100.0	6.0821	5	311	15.2	386.63	29.93	16.5
10	0.17004	12.5	7.87	0	0.5240	6.004	85.9	6.5921	5	311	15.2	386.71	17.10	18.9
11	0.22489	12.5	7.87	0	0.5240	6.377	94.3	6.3467	5	311	15.2	392.52	20.45	15.0
12	0.11747	12.5	7.87	0	0.5240	6.009	82.9	6.2267	5	311	15.2	396.90	13.27	18.9
13	0.09378	12.5	7.87	0	0.5240	5.889	39.0	5.4509	5	311	15.2	390.50	15.71	21.7
14	0.62976	0.0	8.14	0	0.5380	5.949	61.8	4.7075	4	307	21.0	396.90	8.26	20.4
15	0.63796	0.0	8.14	0	0.5380	6.096	84.5	4.4619	4	307	21.0	380.02	10.26	18.2

Dataset: BostonHousing

Label / Target (Y)

	crim	zn	indus	chas	nox	rm	age	dis	rad	tax	ptratio	b	lstat	medv
1	0.00632	18.0	2.31	0	0.5380	6.575	65.2	4.0900	1	296	15.3	396.90	4.98	24.0
2	0.02731	0.0	7.07	0	0.4690	6.421	78.9	4.9671	2	242	17.8	396.90	9.14	21.6
3	0.02729	0.0	7.07	0	0.4690	7.185	61.1	4.9671	2	242	17.8	392.83	4.03	34.7
4	0.03237	0.0	2.18	0	0.4580	6.998	45.8	6.0622	3	222	18.7	394.63	2.94	33.4
5	0.06905	0.0	2.18	0	0.4580	7.147	54.2	6.0622	3	222	18.7	396.90	5.33	36.2
6	0.02985	0.0	2.18	0	0.4580	6.430	58.7	6.0622	3	222	18.7	394.12	5.21	28.7
7	0.08829	12.5	7.87	0	0.5240	6.012	66.6	5.5605	5	311	15.2	395.60	12.43	22.9
8	0.14455	12.5	7.87	0	0.5240	6.172	96.1	5.9505	5	311	15.2	396.90	19.15	27.1
9	0.21124	12.5	7.87	0	0.5240	5.631	100.0	6.0821	5	311	15.2	386.63	29.93	16.5
10	0.17004	12.5	7.87	0	0.5240	6.004	85.9	6.5921	5	311	15.2	386.71	17.10	18.9
11	0.22489	12.5	7.87	0	0.5240	6.377	94.3	6.3467	5	311	15.2	392.52	20.45	15.0
12	0.11747	12.5	7.87	0	0.5240	6.009	82.9	6.2267	5	311	15.2	396.90	13.27	18.9
13	0.09378	12.5	7.87	0	0.5240	5.889	39.0	5.4509	5	311	15.2	390.50	15.71	21.7
14	0.62976	0.0	8.14	0	0.5380	5.949	61.8	4.7075	4	307	21.0	396.90	8.26	20.4
15	0.63796	0.0	8.14	0	0.5380	6.096	84.5	4.4619	4	307	21.0	380.02	10.26	18.2

Dataset: BostonHousing

Features (X)

Label / Target (Y)

	crim	zn	indus	chas	nox	rm	age	dis	rad	tax	ptratio	b	lstat	medv
1	0.00632	18.0	2.31	0	0.5380	6.575	65.2	4.0900	1	296	15.3	396.90	4.98	24.0
2	0.02731	0.0	7.07	0	0.4690	6.421	78.9	4.9671	2	242	17.8	396.90	9.14	21.6
3	0.02729	0.0	7.07	0	0.4690	7.185	61.1	4.9671	2	242	17.8	392.83	4.03	34.7
4	0.03237	0.0	2.18	0	0.4580	6.998	45.8	6.0622	3	222	18.7	394.63	2.94	33.4
5	0.06905	0.0	2.18	0	0.4580	7.147	54.2	6.0622	3	222	18.7	396.90	5.33	36.2
6	0.02985	<h1>Supervised Learning</h1>											5.21	28.7
7	0.08829												12.43	22.9
8	0.14455												19.15	27.1
9	0.21124	12.5	7.87	0	0.5240	5.631	100.0	6.0821	5	311	15.2	386.63	29.93	16.5
10	0.17004	12.5	7.87	0	0.5240	6.004	85.9	6.5921	5	311	15.2	386.71	17.10	18.9
11	0.22489	12.5	7.87	0	0.5240	6.377	94.3	6.3467	5	311	15.2	392.52	20.45	15.0
12	0.11747	12.5	7.87	0	0.5240	6.009	82.9	6.2267	5	311	15.2	396.90	13.27	18.9
13	0.09378	12.5	7.87	0	0.5240	5.889	39.0	5.4509	5	311	15.2	390.50	15.71	21.7
14	0.62976	0.0	8.14	0	0.5380	5.949	61.8	4.7075	4	307	21.0	396.90	8.26	20.4
15	0.63796	0.0	8.14	0	0.5380	6.096	84.5	4.4619	4	307	21.0	380.02	10.26	18.2

Dataset: BostonHousing

Mapping

## Features (X)

	crim	zn	indus	chas	nox	rm	age	dis	rad	tax	ptratio	b	lstat
1	0.00632	18.0	2.31	0	0.5380	6.575	65.2	4.0900	1	296	15.3	396.90	4.98
2	0.02731	0.0	7.07	0	0.4690	6.421	78.9	4.9671	2	242	17.8	396.90	9.14
3	0.02729	0.0	7.07	0	0.4690	7.185	61.1	4.9671	2	242	17.8	392.83	4.03
4	0.03237	0.0	2.18	0	0.4580	6.998	45.8	6.0622	3	222	18.7	394.63	2.94
5	0.06905	0.0	2.18	0	0.4580	7.147	54.2	6.0622	3	222	18.7	396.90	5.33
6	Unsupervised Learning												
7													
8													
9													
10	0.21124	12.5	7.87	0	0.5240	5.631	100.0	6.0821	5	311	15.2	386.63	29.93
11	0.17004	12.5	7.87	0	0.5240	6.004	85.9	6.5921	5	311	15.2	386.71	17.10
12	0.22489	12.5	7.87	0	0.5240	6.377	94.3	6.3467	5	311	15.2	392.52	20.45
13	0.11747	12.5	7.87	0	0.5240	6.009	82.9	6.2267	5	311	15.2	396.90	13.27
14	0.09378	12.5	7.87	0	0.5240	5.889	39.0	5.4509	5	311	15.2	390.50	15.71
15	0.62976	0.0	8.14	0	0.5380	5.949	61.8	4.7075	4	307	21.0	396.90	8.26
16	0.63796	0.0	8.14	0	0.5380	6.096	84.5	4.4619	4	307	21.0	380.02	10.26

Dataset: BostonHousing



## Quick summary

Supervised Learning	Unsupervised Learning
Has features (x) and labels (y)	Has features (x) without labels (y)
The goal is <b>PREDICT</b>	The goal is to <b>SUMMARISE</b>
Example algorithms <ul style="list-style-type: none"><li>- Regression</li><li>- Classification</li></ul>	Example algorithms <ul style="list-style-type: none"><li>- Clustering</li><li>- Association Rules</li><li>- Principal Component Analysis</li></ul>



คอร์สเราโฟกัสที่ supervised learning

AIS อยากจะทำ market survey กับ  
ลูกค้า (ทุกค่าย) ทั้งหมด 3000 คน เพื่อ  
จะดูว่าตลาดคนไทยมีลูกค้าอยู่ที่ประเภท?  
i.e. customer segmentation

Gmail มีตัวกรอง email ว่าอันไหนคือ spam อันไหนคือ ham (อีเมลดี)

R

## Problem 03



อิงเขียนโค้ดทำ web scraping  
จากเว็บไซต์ขายรถยนต์มือสอง  
เพื่อจะดูว่ารถยนต์ Toyota รุ่น  
2015 เครื่อง 1.5 ลิตร ขับมาแล้ว  
20000 โล ควรจะซื้อราคาเท่าไรดี?

น้องอิง!



# Types of Supervised Learning

1. Regression	2. Classification
Predict <b>numeric</b> labels	Predict <b>categorical</b> labels
Examples <ul style="list-style-type: none"><li>- house price</li><li>- customer satisfaction</li><li>- personal income</li><li>- how much a customer will spend</li></ul>	Examples <ul style="list-style-type: none"><li>- yes/ no question</li><li>- churn prediction</li><li>- conversion</li><li>- weather forecast</li><li>- default prediction</li></ul>
100, 200, 250, 190, 300, 500, etc.	0, 1, 0, 0, 1, 1, 0, 0, 1, 1, 0, 0, etc.

R

**ML = Pokemon**







Now let's get  
into the  
details :)


## 3 steps to build ML

- prepare data
- train algorithm
- test/ evaluate algorithm

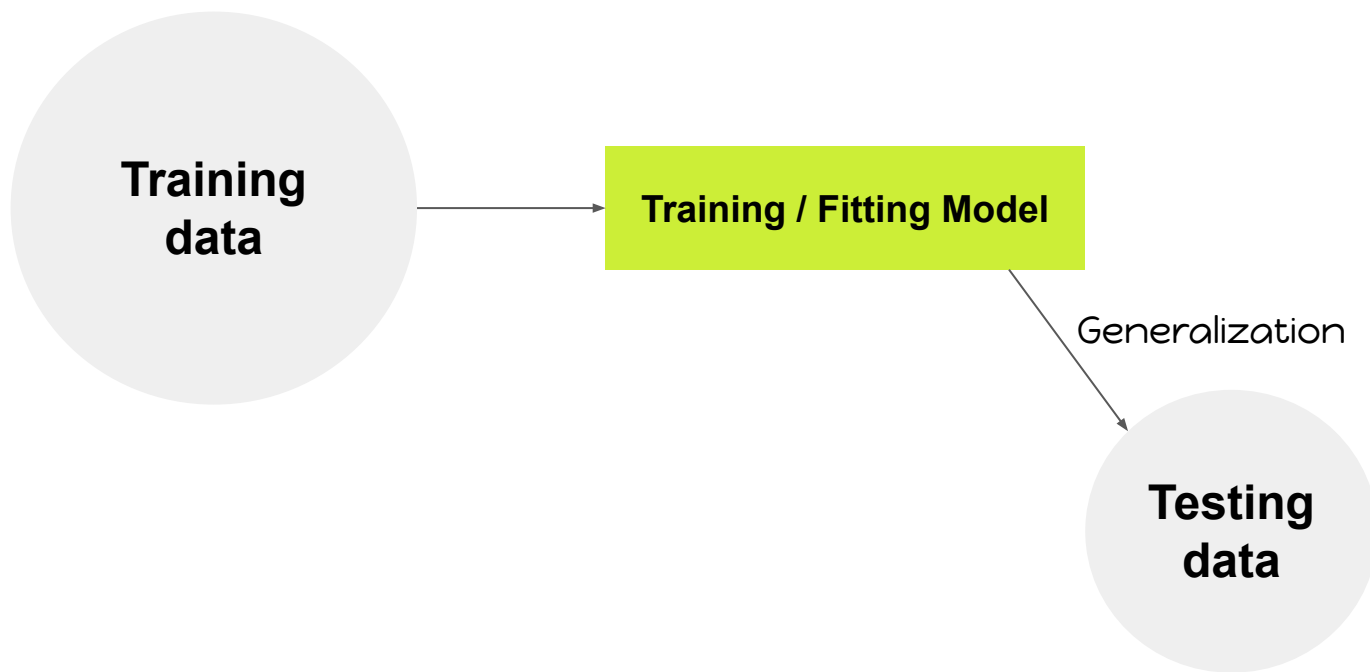


## ML Glossary 02


- train test split
- training set
- testing/ validation set
- overfitting



Train test split

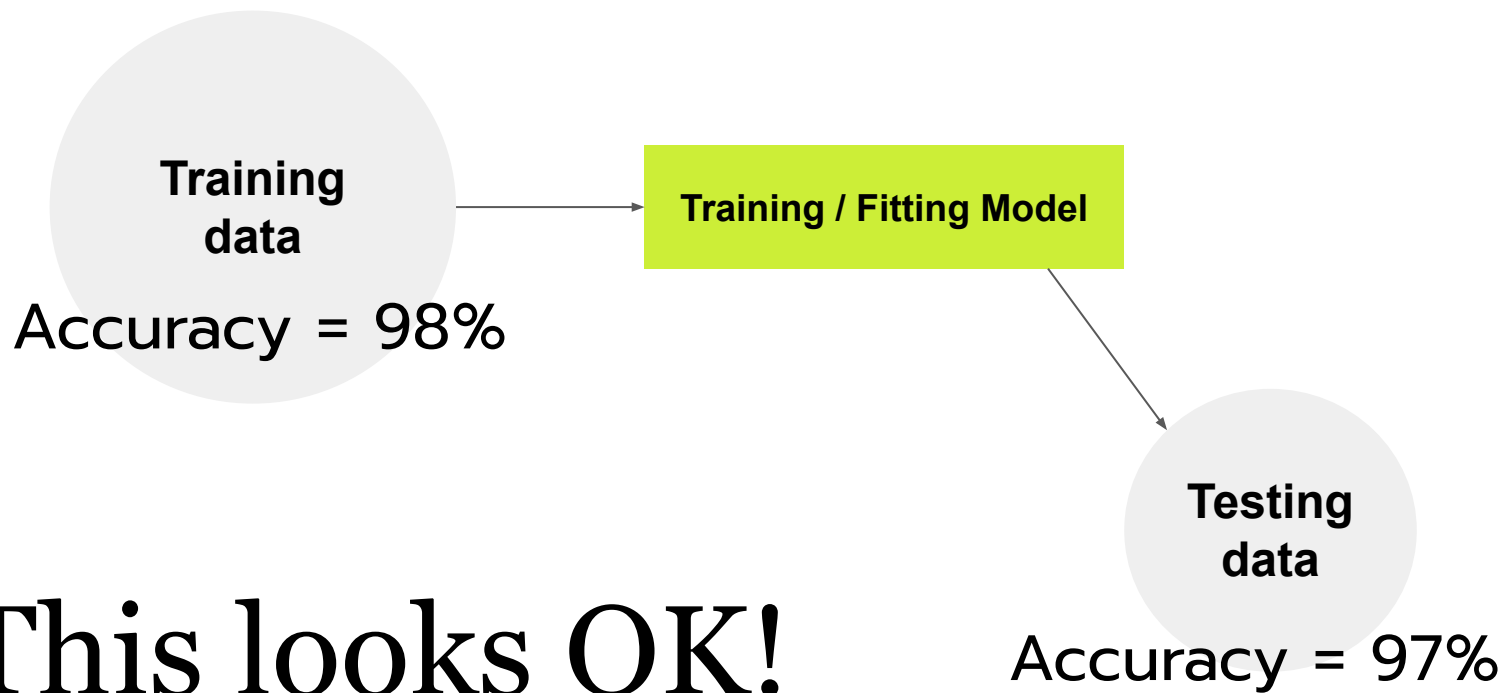


เราต้องตามคำตามนี้เสมอ  
โมเดลที่เราสร้างขึ้นมาเอาไปใช้จริงได้หรือเปล่า? i.e.  
ความถูกต้องของโมเดลกับ test data เป็นเท่าไร



กรณีนี้เรียกว่า **Overfitting**  
โมเดลที่เราสร้างขึ้นมาฟิตกับข้อมูล  
Training มากเกินไปจนไม่สามารถนำไป  
ใช้กับ Testing/ Unseen data ได้






**This looks OK!**

## Golden Rule


เราจะไม่ทดสอบโมเดลด้วยข้อมูลชุดเดิมที่ใช้เทรนโมเดล

i.e. เราจะไม่ใช้ training data วัดผลว่าโมเดลของเรากำลังทำงานดีไหม? แต่ต้องเป็น unseen data ที่โมเดลไม่เคยเห็นมาก่อน


## R Our goal is in the middle -> Just Right



Underfitted



Good Fit/Robust



Overfitted

<https://medium.com/greyatom/what-is-underfitting-and-overfitting-in-machine-learning-and-how-to-deal-with-it-6803a989c76>

# **Discuss:** Overfitting คืออะไร?

## เขียนคำตอบได้ที่นี้

# **Discuss:** แล้วถ้า Underfitting ล่ะ?

## เขียนคำตอบได้ที่นี้

ในทางปฏิบัติ Train Test Split (ส่วนมาก) จะไม่  
ใช้วิธีที่ดีที่สุดในการสร้างโมเดล ML

เราใช้เทคนิคที่เรียกว่า **Resampling** สำหรับเทรน  
โมเดลเพื่อผลลัพธ์ที่ดีกว่า




- resampling
  - leave one out CV
  - bootstrap
  - k-fold cross validation






## Leave One Out CV





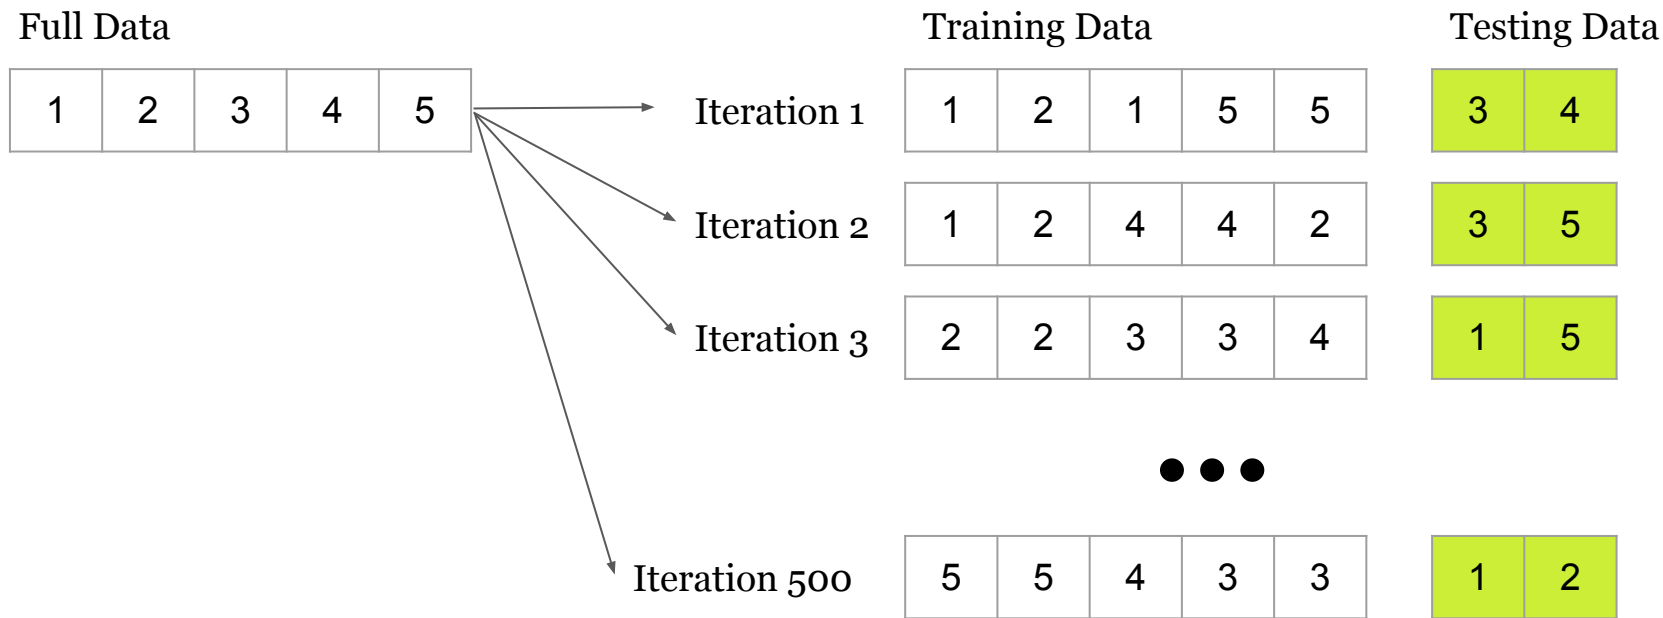
## Leave One Out CV

1	2	3	4	...	...	997	998	999	1000	iteration 1
1	2	3	4	...	...	997	998	999	1000	iteration 2
1	2	3	4	...	...	997	998	999	1000	iteration 3
1	2	3	4	...	...	997	998	999	1000	iteration 4
1	2	3	4	...	...	997	998	999	1000	iteration 5
■ ■ ■										
1	2	3	4	...	...	997	998	999	1000	iteration 999
1	2	3	4	...	...	997	998	999	1000	iteration 1000



Sampling with replacement  
ใช้การสุ่มซ้ำ  $n=1000$  เหมือน full dataset

# R Bootstrap



The error will be averaged over 500 training iterations



## K-Fold Cross Validation

1	2	3	4	5
1	2	3	4	5
1	2	3	4	5
1	2	3	4	5
1	2	3	4	5

iteration 1: train {2,3,4,5} test {1} -> error 18%

iteration 2: train {1,3,4,5} test {2} -> error 20%

iteration 3: train {1,2,4,5} test {3} -> error 30%

iteration 4: train {1,2,3,5} test {4} -> error 15%

iteration 5: train {1,2,3,4} test {5} -> error 19%

Average error =  $(18+20+30+15+19) / 5 = 20.4\%$

ปกติเรานิยมใช้ค่า **K=5** หรือ **K=10**

# **Discuss:** LOOCV, Bootstrap, K-Fold ทั้งสามวิธีแตกต่างกันอย่างไร?

## เขียนคำตอบได้ที่นี้

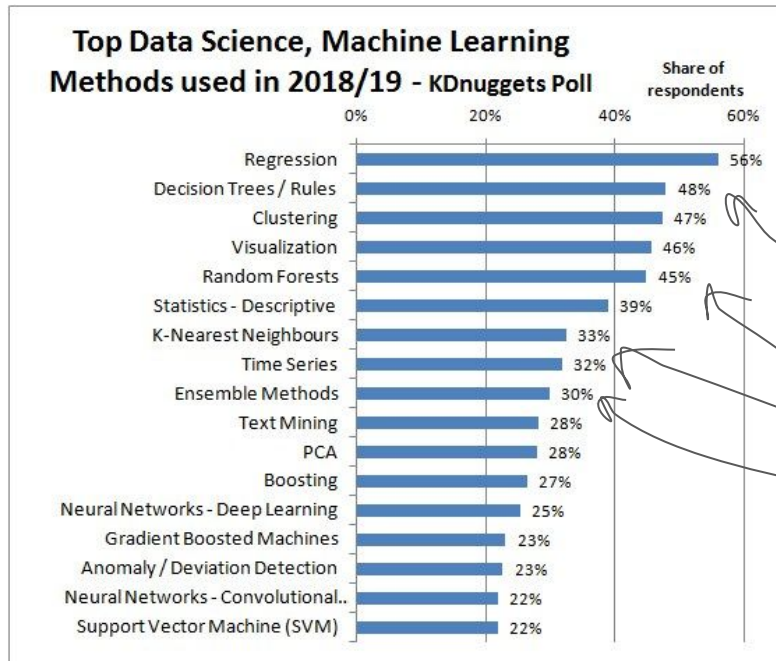


## Essential ML

- OK** what exactly is machine learning
- OK** supervised vs. unsupervised
- OK** regression vs. classification
- OK** train test split vs. cross validation
  - model selection + hyperparameter
  - model evaluation



# Popularity of Algorithms



Today we'll cover

- Linear Regression
- Logistic Regression
- Regularized Regression
- Decision Tree
- Random Forests (Bagging)
- K-Nearest Neighbours
- Ensemble Model

<https://www.kdnuggets.com/2019/04/top-data-science-machine-learning-methods-2018-2019.html>





## No Free Lunch

No Free Lunch แปลว่า “ไม่มีโมเดลไหนเก่งที่สุด และสามารถตอบโจทย์ได้ทุกปัญหา”

ถ้ามีใครถามว่าโมเดลไหนเก่งที่สุด?

ให้ตอบว่า “It depends” (ขึ้นอยู่กับข้อมูล)

ความท้าทายของ ML คือการหาโมเดลที่ดีที่สุดสำหรับปัญหาที่เรากำลังแก้

## R Occam's Razor

Algorithm #1

vs.

Algorithm #2

ถ้ามีโมเดลสองตัวที่มี performance ดีเท่าๆกัน ให้เลือกตัวที่สร้างและอธิบายได้ง่ายกว่า (**choose simpler model**)



## How to choose a model


ให้ลองถาม 2 คำถามง่ายๆนี้

1. ปัญหานี้เป็น regression หรือ classification?
2. อยากได้ high accuracy หรือ high interpretability?
  - Always choose a simpler model if performances are similar
  - Try different algorithms and find the right one.



# Caret Package

Learn more at <https://topepo.github.io/caret/index.html>



**Max Kuhn**

the author of caret package





## Dataset for our projects

```
## load library
## install.packages("mlbench")
library(mlbench)
library(tidyverse)

## load dataset for regression
data("BostonHousing")
glimpse(BostonHousing)

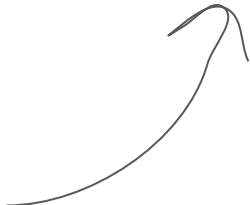
## load dataset for classification
data("PimaIndiansDiabetes")
glimpse(PimaIndiansDiabetes)
```





## Caret Training Template

```
model <- train(form = y ~ . ,  
               data = train_data ,  
               method = "lm" )
```

Model that we  
want to train




# R Our first machine



## R Euclidean Distance

$$d = \sqrt{(x_1 - x_2)^2 + (y_1 - y_2)^2}$$






## Euclidean Distance

$$d = \sqrt{(x1 - x2)^2 + (y1 - y2)^2}$$


```
point_1 <- c(2,3)
point_2 <- c(6,8)
d <- sqrt( (2-6)**2 + (3-8)**2 )
print(d)
```



## We use majority vote to assign label



Predict 'Red'  $3/5 = 60\%$



Predict 'Green'  $3/5 = 60\%$

## Majority Vote




### Steps to train this algorithm

1. แต่งตั้ง สว 250 คน
2. ช่วยกันเลือกนายกฯ

เฮ้ย เต๋วๆๆๆ 555555555555+

## **R** We use average value for regression problem



Use the average as prediction

$$(10 + 16 + 18 + 20 + 15) / 5 = 15.8$$



## **K-Nearest Neighbors**

1. Choose K
2. Compute distance
3. Majority vote for classification or  
Average for regression

## Train test split (the easiest method)

Prepare dataset first

We'll use split data into training 75% and testing 25%

```
## split data
set.seed(99)
n <- nrow(BostonHousing)
id <- sample(n, size = n*0.75, replace=FALSE)
train_data <- BostonHousing[id, ]
test_data <- BostonHousing[-id, ]
```





## Very easy to train a machine in R

```
## train model
set.seed(99)
knn_model <- train(medv ~ .,
                   data = train_data,
                   method = "knn")

## test model
p <- predict(knn_model, newdata = test_data)

## rmse
rmse <- sqrt(mean((p - test_data$medv)**2))
```





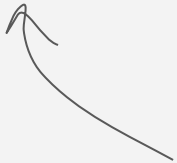
## KNN + K-Fold CV

```
## train model
set.seed(99)
ctrl <- trainControl(method = "cv", number = 5, verboseIter = TRUE)
knn_model <- train(medv ~ .,
                   data = train_data,
                   method = "knn",
                   trControl = ctrl)

## test model
p <- predict(knn_model, newdata = test_data)

## rmse
rmse <- sqrt(mean((p - test_data$medv)**2))
```

5 Fold Cross Validation

A hand-drawn arrow points from the text "5 Fold Cross Validation" to the number "5" in the `number = 5` parameter of the `trainControl` function.



Tuning



## Random Search

```
## train model
set.seed(99)
ctrl <- trainControl(method = "cv", number = 5, verboseIter = TRUE)
knn_model <- train(medv ~ .,
                  data = train_data,
                  tuneLength = 5,
                  method = "knn",
                  trControl = ctrl)
```

Try 5 values of K



```
## test model
p <- predict(knn_model, newdata = test_data)

## rmse
rmse <- sqrt(mean((p - test_data$medv)**2))
```



## Grid Search

```
## create grid
myGrid <- expand.grid(k = 1:10)

## train model
set.seed(99)
ctrl <- trainControl(method = "cv", number = 5, verboseIter = TRUE)
knn_model <- train(medv ~ .,
  data = train_data,
  tuneGrid = myGrid,
  method = "knn",
  trControl = ctrl)

## test model
p <- predict(knn_model, newdata = test_data)

## rmse
rmse <- sqrt(mean((p - test_data$medv)**2))
```

The values we  
select ourselves :D



R

## Grid Search Result

```
> knn_model
k-Nearest Neighbors

379 samples
13 predictor

No pre-processing
Resampling: Cross-Validated (5 fold)
Summary of sample sizes: 303, 303, 303, 303, 304
Resampling results across tuning parameters:
```

k	RMSE	Rsquared	MAE
1	7.607647	0.4598462	4.922035
2	6.857741	0.5208036	4.750591
3	6.778822	0.5139722	4.657967
4	6.659557	0.5153593	4.651180
5	6.646851	0.5131619	4.681624
6	6.660705	0.5081547	4.648119
7	6.711653	0.5001402	4.661983
8	6.881981	0.4749499	4.749852
9	6.872293	0.4768345	4.763948
10	6.927021	0.4683690	4.773996

RMSE was used to select the optimal model using the smallest value.  
The final value used for the model was k = 5.

Cross Validation  
ช่วยเราเลือกค่า k ที่  
ทำให้ RMSE ต่ำที่สุด  
ตอนเรา train model

1. KNN เข้าใจง่ายทำงานได้โอเคร ถ้า feature ไม่เยอะมาก
2. KNN ใช้ได้ทั้ง regression/ classification
3. K ใน KNN คือค่า hyperparameter ที่เราเปลี่ยนได้
4. เราเลือก K ที่ทำให้ train RMSE ต่ำที่สุด
5. train RMSE ต่ำที่สุดไม่ได้แปลว่าโมเดลเราจะทำนาย test\_data ได้ดี ต้องเอาไปทดสอบอีกที





# **Caret Interface Summary**



# Classification vs. Regression

## Classification

```
set.seed(42)

ctrl <- trainControl(method = "cv",
                     number = 5)

model <- train(
  y ~ .,
  data = df,
  method = "knn",
  metric = "Accuracy",
  trControl = ctrl
)
```


## Regression

```
set.seed(42)

ctrl <- trainControl(method = "cv",
                     number = 5)

model <- train(
  y ~ .,
  data = df,
  method = "knn",
  metric = "RMSE",
  trControl = ctrl
)
```

Same interface, different metrics





# Classification Interfaces

## Classification - ROC Sens Specs

```
set.seed(42)

ctrl <- trainControl(
  method = "cv",
  number = 5,
  summaryFunction = twoClassSummary,
  classProbs = TRUE)

model <- train(
  y ~ .,
  data = df,
  method = "knn",
  metric = "ROC",
  trControl = ctrl
)
```

## Classification - AUC Precision Recall F1



```
set.seed(42)

ctrl <- trainControl(
  method = "cv",
  number = 5,
  summaryFunction = prSummary,
  classProbs = TRUE)

model <- train(
  y ~ .,
  data = df,
  method = "knn",
  metric = "AUC",
  trControl = ctrl
)
```



# Linear Regression Explained




## Linear Regression Model

$$\hat{y} = B_0 + B_1 X_1$$

prediction

y intercept

slope



Linear Regression finds  
B0 and B1 that minimize  
the error

## Minimize Error

$$\text{minimize } \sum (\text{prediction} - \text{actual})^2$$

$$\text{minimize } \sum (\hat{y} - y)^2$$



Sum of Squared Error  
or RSS (for short)

## R

## Common Regression Metrics

$$MAE = \frac{1}{n} * \sum |\hat{y} - y|$$

$$MSE = \frac{1}{n} * \sum (\hat{y} - y)^2$$

$$RMSE = \sqrt{\frac{1}{n} * \sum (\hat{y} - y)^2}$$

โมเดลที่เราเทรนจะพยายามทำให้ค่า  
MAE/ MSE/ RMSE มีค่าต่ำที่สุด  
i.e. minimize error



## Easy to compute in Spreadsheets

y	y_hat	error	error	error^2	
10	8.5	1.5	1.5	2.25	
12	14.5	-2.5	2.5	6.25	
14	10	4	4	16	
16	17	-1	1	1	
18	17.5	0.5	0.5	0.25	
			<b>9.5</b>	<b>25.75</b>	
			1.9	5.2	2.3
			MAE	MSE	RMSE



## Build linear regression in R

```
## train model with train_data
set.seed(99)
lm_model <- train(medv ~ rm + indus + crim,
                  data = train_data,
                  method = "lm")

## test model (predict test data)
p <- predict(lm_model, newdata = test_data)
rmse <- sqrt(mean( (p - test_data$medv)** 2 ))
```





## Linear regression with K-Fold

```
## train model with train_data
set.seed(99)
ctrl <- trainControl(method = "cv", number = 5,
                     verboseIter = TRUE)

lm_model <- train(medv ~ rm + indus + crim,
                  data = train_data,
                  method = "lm",
                  trControl = ctrl)

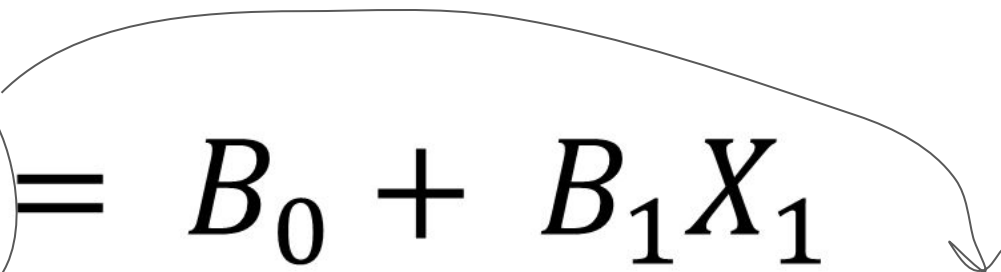
## test model (predict test data)
p <- predict(lm_model, newdata = test_data)
rmse <- sqrt(mean( (p - test_data$medv)** 2 ))
```

# R Logistic regression for binary classification





**R** Logistic is very similar to linear regression

$$Z = B_0 + B_1 X_1$$



$$P(Y = 1|x) = \frac{e^z}{1 + e^z}$$

Sigmoid function



Z	sigmoid(Z)	y_hat
-10	0.0000	0
-9	0.0001	0
-8	0.0003	0
-7	0.0009	0
-6	0.0025	0
-5	0.0067	0
-4	0.0180	0
-3	0.0474	0
-2	0.1192	0
-1	0.2689	0
0	0.5000	0
1	0.7311	1
2	0.8808	1
3	0.9526	1
4	0.9820	1
5	0.9933	1
6	0.9975	1
7	0.9991	1
8	0.9997	1
9	0.9999	1
10	1.0000	1



If  $\text{sigmoid}(Z) > 0.5$ , predict  $y = 1$ , else  $y = 0$



## R Our prediction changes if threshold changes



## R Code

```
## train model with train_data
set.seed(99)
ctrl <- trainControl(method = "cv", number = 5,
                      verboseIter = TRUE)

logistic_model <- train(diabetes ~ .,
                        data = train_data,
                        method = "glm",
                        trControl = ctrl)

## test model (predict test data)
p <- predict(logistic_model, newdata = test_data)
accuracy <- mean(p == test_data$diabetes)
```



- Accuracy
- Precision
- Recall
- F1




สามารถคำนวณได้  
ง่ายๆจาก

**Confusion Matrix**

R

## Common classification metrics



AUC  
Area Under Curve



## Interpretation in Thai

Metrics	ความหมาย
Accuracy	ความถูกต้องของโมเดลในภาพรวม
Precision	ทุก 100 ครั้งที่เรากำหนด $y=1$ โอกาสถูกเท่าไร
Recall	ทุกผู้ป่วยจริงๆ 100 คน เราตรวจเจอที่คน
F1	ค่าเฉลี่ยระหว่าง precision, recall

### 27.4.5 Balanced accuracy and $F_1$ score

Although we usually recommend studying both specificity and sensitivity, very often it is useful to have a one-number summary, for example for optimization purposes. One metric that is preferred over overall accuracy is the average of specificity and sensitivity, referred to as *balanced accuracy*. Because specificity and sensitivity are rates, it is more appropriate to compute the *harmonic average*. In fact, the  $F_1$ -score, a widely used one-number summary, is the harmonic average of precision and recall:

$$\frac{1}{\frac{1}{2} \left( \frac{1}{\text{recall}} + \frac{1}{\text{precision}} \right)}$$

Because it is easier to write, you often see this harmonic average rewritten as:

$$2 \times \frac{\text{precision} \cdot \text{recall}}{\text{precision} + \text{recall}}$$

when defining  $F_1$ .





## R Code - Confusion Matrix

```
## use table()
table(predicted, actual, dnn = c("predicted", "actual"))
```

		Actual	
		neg	pos
Predicted	neg	101 <b>TN</b>	33 <b>FN</b>
	pos	14 <b>FP</b>	44 <b>TP</b>

```
## how we calculate four metrics
accuracy <- (101 + 44) / (101 + 33 + 44 + 14)
precision <- 44 / (44 + 14)
recall <- 44 / (44 + 33)
F1 <- 2 * (precision * recall) / (precision + recall)
```



## Regularized Regression

1. Ridge Regression
2. Lasso Regression

Regularization is a key technique in ML to **reduce overfitting** :D

R

## Lasso Regression (L1)

$$RSS = \sum (\hat{y} - y)^2$$

Normal RSS from Linear Regression

$$Lasso\ RSS = \sum (\hat{y} - y)^2 + \lambda \sum |\beta|$$

Lasso add this term to the error function

R

## Ridge Regression (L2)

$$RSS = \sum (\hat{y} - y)^2$$

Normal RSS from Linear Regression

$$Ridge\ RSS = \sum (\hat{y} - y)^2 + \lambda \sum \beta^2$$

Ridge add this term to the error function

## **R** Regularization helps reduce overfitting

$$y_{\text{hat}} = b_0 + b_1x_1 + b_2x_2 + b_3x_3 + b_4x_4$$

$$y_{\text{hat}} = 100 + \mathbf{150}x_1 + \mathbf{200}x_2 + \mathbf{120}x_3 + \mathbf{80}x_4$$


$$\text{Lasso} = \text{RSS} + \text{Lambda} * (150 + 200 + 120 + 80)$$

Lambda = 0

Error is the same as Linear Regression

Lambda > 0

All coefficient (B) in the model must be **shrunk** to reduce the new error

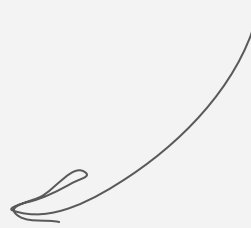
script.R

```
1 # Train glmnet with custom trainControl and tuning: model
2 model <- train(
3   y ~ .,
4   data = overfit,
5   tuneGrid = expand.grid(
6     alpha = 0:1,
7     lambda = seq(0.0001, 1, length=20)
8   ),
9   method = "glmnet",
10  trControl = myControl
11 )
12
13 # Print model to console
14 model
15
16 # Print maximum ROC statistic
17 max(model$results$ROC)
```

```
## train elasticnet model
set.seed(99)
ctrl <- trainControl(method = "cv", number = 5,
                     verboseIter = TRUE)

enet_model <- train(diabetes ~ .,
                   data = train_data,
                   method = "glmnet",
                   trControl = ctrl)
```

ElasticNet =  
Mixed between  
Ridge + Lasso



```
## test model
p <- predict(enet_model, newdata = test_data)
accuracy <- mean(p == test_data$diabetes)
```





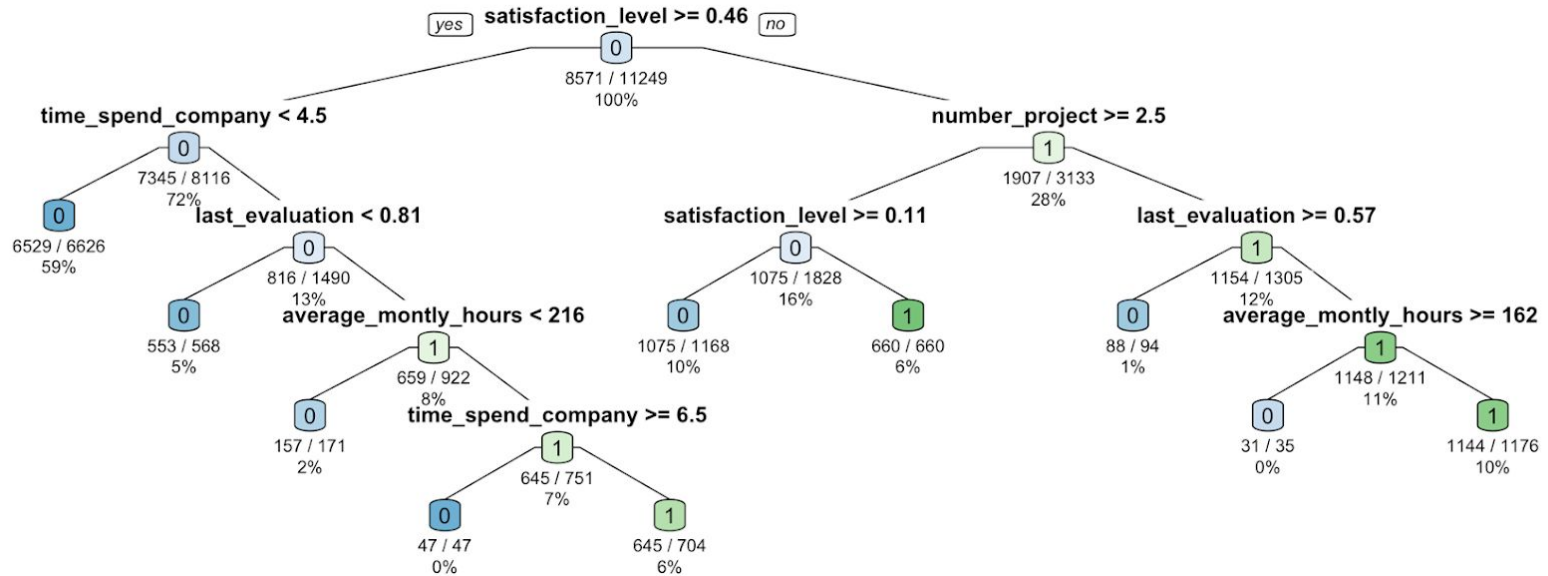
**Time for a fun game :D**

- Ask me a yes/no question
- To guess my favourite animal
- Max 10 questions











# Decision Tree





## How decision tree work?







Age	Sex	App
15	M	 tinder
20	M	 tinder
16	F	
19	M	 tinder
22	F	
20	F	

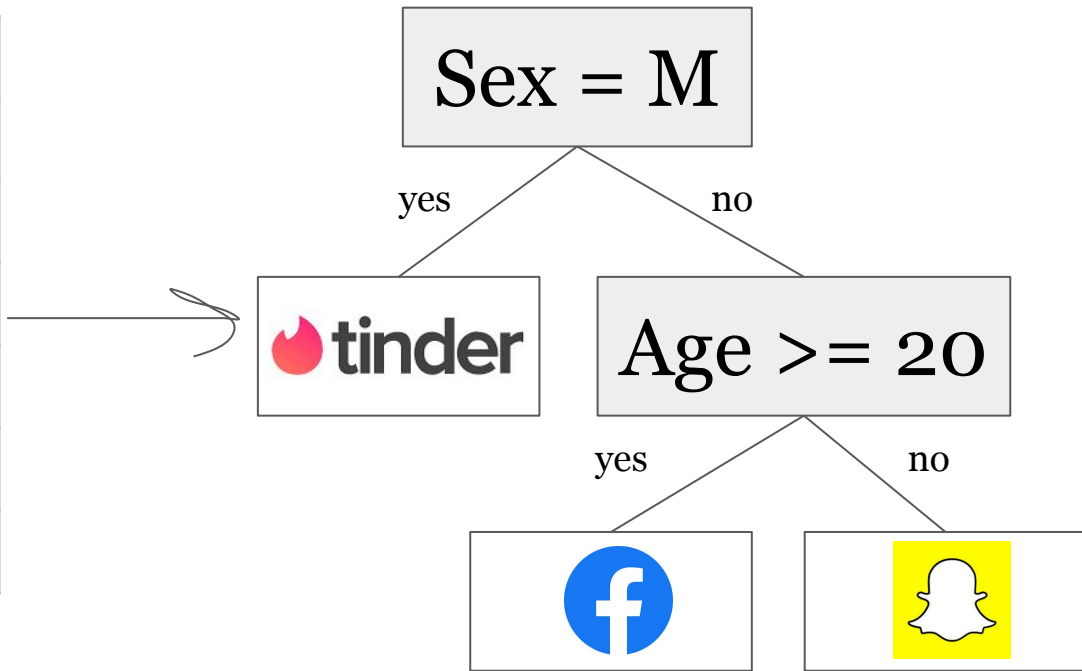
เวลาเราสร้าง decision tree เราถามคำถาม yes/no question ทีละข้อ

i.e. feature ใช้แบ่ง App ได้ดีที่สุด



## How decision tree work?

Age	Sex	App
15	M	 tinder
20	M	 tinder
16	F	
19	M	 tinder
22	F	
20	F	



## Decision Tree with K-Fold

```
## train tree
set.seed(99)

ctrl <- trainControl(method = "...", number = ...,
                     verboseIter = TRUE)

tree_model <- train(diabetes ~ .,
                   data = ...,
                   method = "rpart",
                   trControl = ...)

## test model
p <- predict(tree_model, newdata = ...)
accuracy <- mean(...)
```



## Let's do a quick review

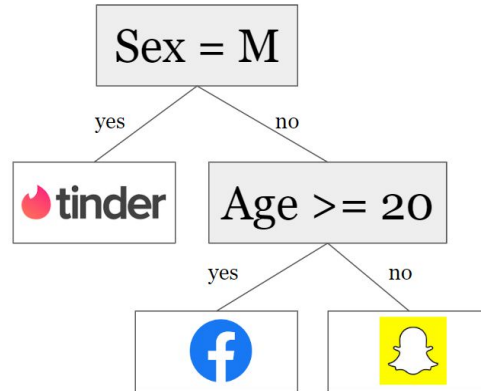

<b>Parametric</b>	<b>Non-Parametric</b>
Linear Regression	KNN
Logistic Regression	Decision Tree
Ridge Regression	Random Forest
Lasso Regression	

Regression is a Linear Combination (Parametric)

$$\hat{y} = b_0 + b_1x_1 + b_2x_2 + b_3x_3 + b_4x_4$$

↖  
มีพจน์

While a non-parametric has no form :P





## Random Forest



We grow hundreds of  
uncorrelated (decision) trees





Combine them to make  
prediction (similar to KNN,  
majority vote or average)


# R Teamwork (Bagging)

## Bootstrap + mtry hyperparameter


All features




Tree 1



Tree 2



Tree 3



Tree 100





## Random Forest Code

```
## train random forests
set.seed(99)
myGrid <- expand.grid(mtry = 2:4)

ctrl <- trainControl(method = "...", number = ...,
                     verboseIter = TRUE)

rf_model <- train(diabetes ~ .,
                  data = ...,
                  method = "rf",
                  tuneGrid = myGrid,
                  trControl = ...)

## test model
p <- predict(rf_model, newdata = ...)
accuracy <- mean(...)
```






## Ensemble Models

Ensemble

American pronunciation ▾

aan · **saam** · bl 🔊

☐ Slow



Feedback

อาน ซาม เบิ้ล



# นำโมเดลหลายๆตัวมาช่วยกัน ทำนายผล (Majority Vote)

KNN	Logistic Regression	Ridge Regression	Decision Tree	Random Forest
1	0	0	1	1



## **Save our models for later use**

```
saveRDS(model, "model.rds")  
  
model <- readRDS("model.rds")
```

- Machine Learning is **art + science**
- Try different models (No Free Lunch)
- Choose the simpler model
- Use CV + Grid Search to fine-tune model
- Start with Regression or decision tree because **they are very fast to train**

# Bootcamp Live 10

## Introduction to Machine Learning

Website: <https://datarockie.com>

